

Rebuttal Material for Paper # 4515

This document contains the following items:

- A new modified Assurance game and new MANSA CL call plot (Table 1 and Figure 2)
- Bar chart of returns for MANSA and baselines for various α values (Figure 1)
- Area under curve results for StarCraft Multi-Agent Challenge (Figure 3 and Table 2)
- Area under curve results for Level-Based Foraging (Figure 4 and Table 3)
- MANSA SMAC maps with extra seeds Figure 5
- Win rate training curves for MANSA with restriction on CL updates and baselines in SMAC (Figure 7)
- Return training curves for MANSA with restriction on CL updates and baselines in LBF (Figure 7)
- Win rates for MANSA-B (MANSA with budget constraint) with restriction on CL updates and baselines (Table 4)
- Pseudocode for MANSA (Figure 8)
- Diagram of MANSA schematic

	Up	Down
Up	$5(1 + \alpha), 5(1 + \alpha)$	$10\alpha, 10\alpha$
Down	$10\alpha, 10\alpha$	10, 10

Table 1: Modified reward functions of Assurance Game.

	AUC	std
MANSA	0.667676	0.013114
QMIX	0.621228	0.048804
IQL	0.3949	0.009503

Table 2: Mean and standard deviation of each algorithm at each step in each map from SMAC.

MANSA with CL training restriction (MANSA_CLR) imposes the restriction that CL updates can only occur when the Global agent makes a CL call (i.e. when $g = 1$). The following training plots display the results for MANSA_CLR against the baselines.

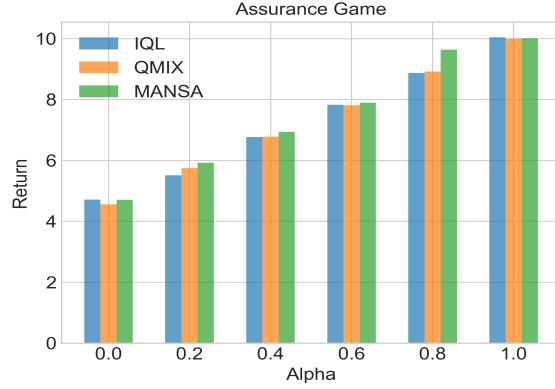


Figure 1: **Normal form games.** End-of-training returns for MANSA, QMIX and IQL for various values of α in the Assurance Game. Despite MANSA reducing its dependence on CL as $\alpha \rightarrow 1$, it achieves returns that are better or the same as the baselines for all α .




Figure 2: MANSA CL calls in Modified Assurance game in Table 1.

	AUC	std
MANSA	0.692759	0.007762
QMIX	0.363176	0.009317
IQL	0.617038	0.006738

Table 3: Mean and standard deviation of each algorithm at each step in each map from LBF.

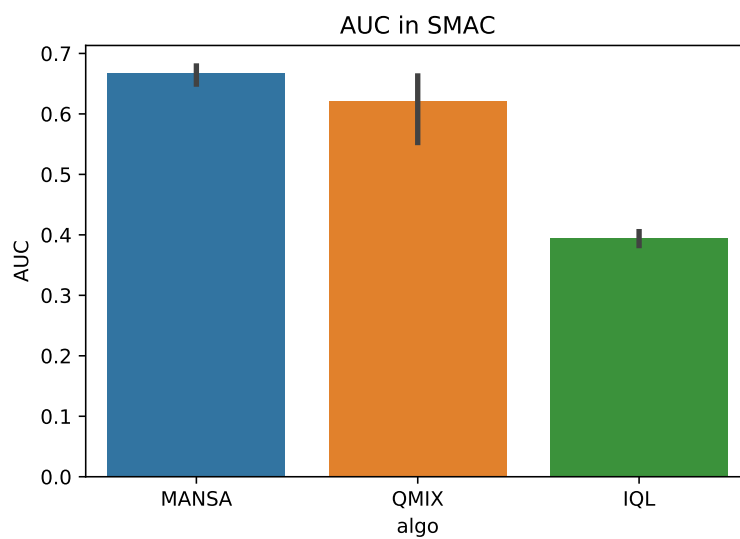


Figure 3: Area under the curve (normalised) results in all tested StarCraft Multi-agent Challenge maps.

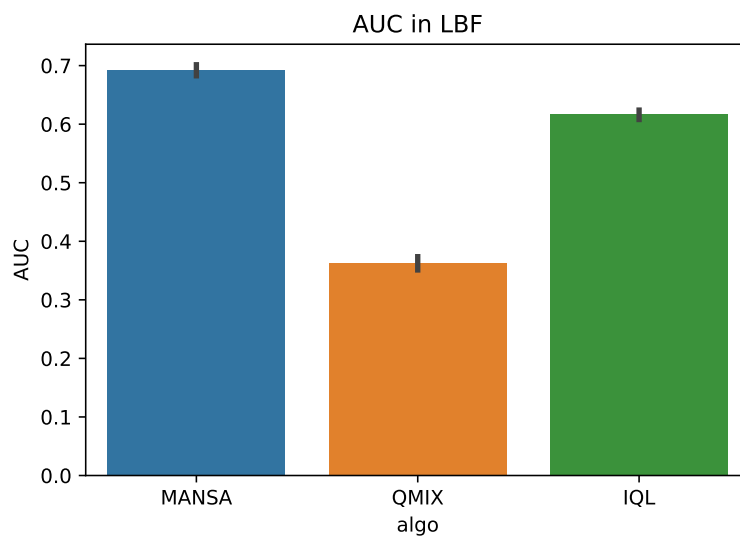


Figure 4: Area under the curve (normalised) results in all tested Level-Based Foraging maps.




Figure 5: MANSA and baselines with 5 seeds in the StarCraft Multi-Agent Challenge (SMAC) map *Corridor*.

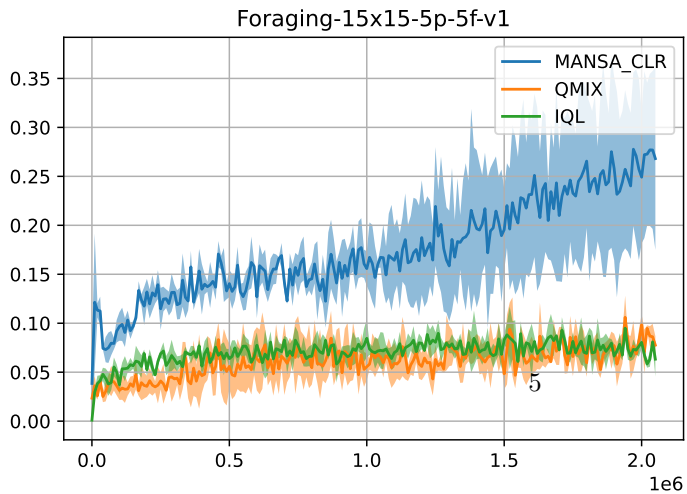




Figure 6: End-of-training returns of MANSA with implementation with CL update restriction (MANSA_CLR) in Level-Based Foraging (LBF).






Figure 7: End-of-training win-rates of MANSA with implementation with CL update restriction (MANSA_CLR) in StarCraft Multi-Agent Challenge (SMAC).

	Original/QMIX/IQL	10%	20%	50%	75%
2m_vs_1z	98.00 ± 1.00 92.00 ± 1.63 87.00 ± 0.82	100.00 ± 0.00	99.67 ± 0.57	96.67 ± 3.05	99.00 ± 0.00
2s3z	96.67 ± 1.24 95.67 ± 1.8 79.67 ± 6.69	82.00 ± 1.41	82.33 ± 5.18	81.67 ± 1.69	96.33 ± 0.47

Table 4: End-of-training win-rates of MANSA-B with implementation with CL update restriction and various CL call budget constraints against baselines.

Algorithm 1 Multi Agent Network Selection Algorithm (MANSA)

Input: Independent policies π^i , centralised policies π^c , Global policy g_0 , independent learning algorithm Δ^i , centralised learning algorithm Δ^c , learning algorithm for Global Δ^g , experience buffer B

Output: Optimised policies π^{i^*} , π^{c^*} , and g^*

for $t = 1, T$ **do**

Given environment state s_t evaluate $g_t \sim g(\cdot|s_t)$

if $g_t = 1$ **then**

Sample action using global state $a_t \sim \pi^c(\cdot|s_t)$ **Use**
Central

else

Sample action using local observations $a_t \sim$
 $\pi^d(\cdot|\tau_t)$ **Use** Independent

Apply action a_t to environment to obtain s_{t+1} , τ_{t+1}
and $r_{t+1} := \sum_{i \in \mathcal{N}} r_{i,t+1}$

Store $(s_t, \tau_t, a_t, r_{t+1}, s_{t+1}, \tau_{t+1})$ in B

if $g_t = 1$ **then**

Sample B to obtain (s_i, a_i, r_i, s_{i+1}) and update π^c
with Δ^c (**Discard** τ_t, τ_{t+1})

else

Sample B to obtain $(\tau_i, a_i, r_i, \tau_{i+1})$ and update
 π^i with Δ^i (**Discard** s_t, s_{t+1})

Sample B to obtain (s_i, g_i, r_i, s_{i+1}) and update g with
 Δ^g (**Discard** a_t, τ_t, τ_{t+1})

Figure 8: Pseudocode for MANSA. This includes a centralised learning update restriction.

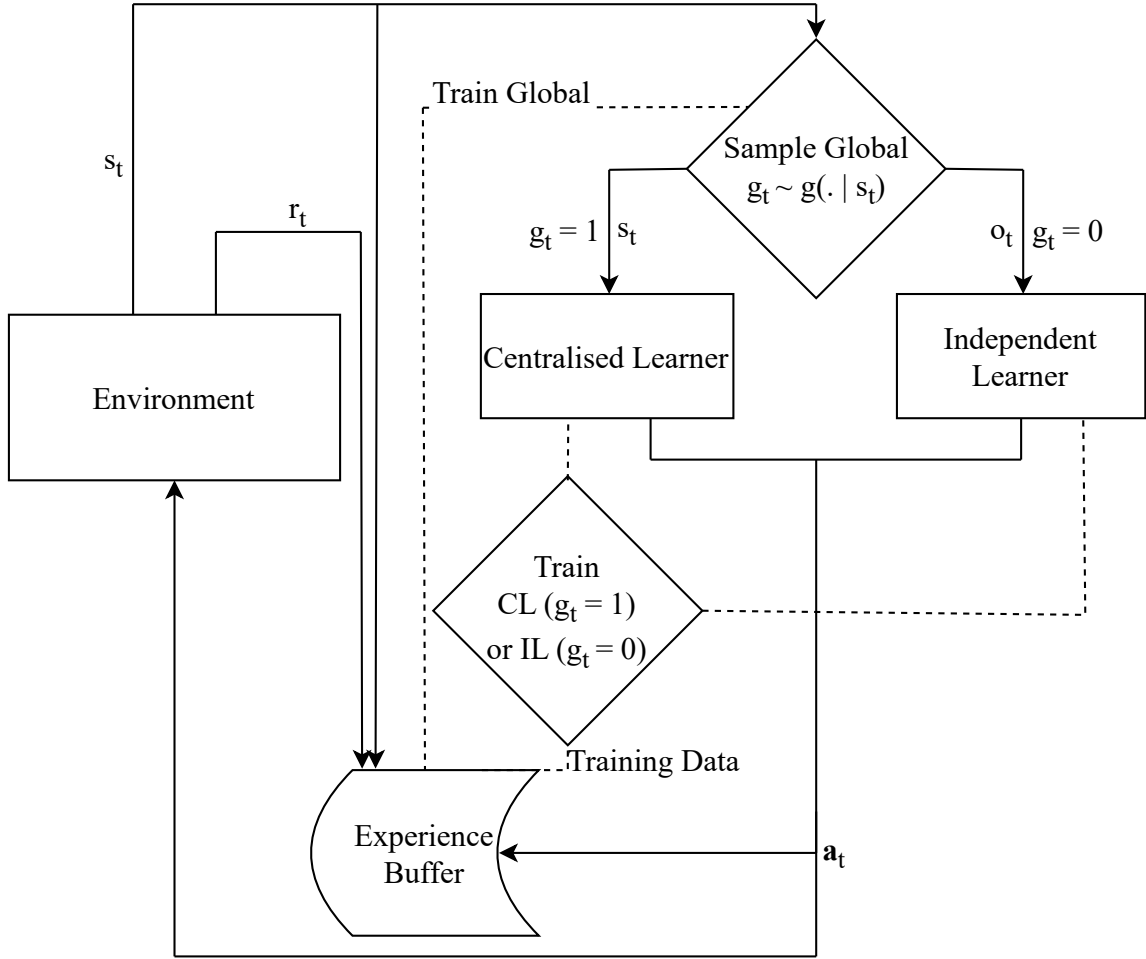


Figure 9: MANSA schematic highlighting the CL update restriction.